

Aryan Yadav

AI/ML & Generative AI Developer

ay6033756@gmail.com | +91 8318695055 | [LinkedIn](#) | [Github](#) | [Leetcode](#) | [Portfolio](#)

Education

IEC College of Engineering & Technology, Greater Noida, India

Sep 2023 – Present

B.Tech – Computer Science Engineering

– Expected Graduation: Jul 2027

Vishnu Bhagwan Public School, Prayagraj, India

2020 – 2023

12th Standard (2022–23) • 10th Standard (2020–21)

Technical Skills

Languages: Python (primary), SQL

AI / ML: Scikit-learn, TensorFlow, Keras, ANN, CNN, RNN, LSTM/GRU

GenAI & LLMs: LangChain, Hugging Face, Transformers, RAG, Embeddings, Prompt Engineering, LLM APIs

Vector Databases: FAISS, ChromaDB, Qdrant

Computer Vision: OpenCV, YOLOv8, MediaPipe

NLP: TF-IDF, NLTK, Tokenization, Text Classification

Data: Pandas, NumPy, Matplotlib, Feature Engineering, Data Preprocessing

Deployment & Tools: Streamlit, FastAPI, Vercel, Docker, Git, GitHub, Supabase, Render, MCP

Projects

DocuMind AI– AI Document Intelligence Platform

Next.js, FastAPI, Python, Qdrant, Supabase, Gemini, Groq | [GitHub Link](#) | [Live Demo](#)

- Built a full-stack AI document intelligence platform supporting PDF, DOCX, and image uploads with semantic search and citation-grounded conversational QA using Next.js, FastAPI, Supabase, and Qdrant.
- Engineered a multimodal RAG pipeline using Docling/OCR, structure-aware chunking, SentenceTransformers embeddings, Qdrant retrieval, and Gemini vision enrichment for context-aware document understanding.
- Implemented Groq/Gemini model fallback, SSE token streaming, authentication with PostgreSQL RLS, Langfuse observability, rollback handling, and automated backend testing.

Insurance Premium Prediction API

Python, FastAPI, Pydantic, scikit-learn, Docker | [GitHub Link](#)

- Built a FastAPI microservice that classifies insurance premium categories (Low/Medium/High) from user demographic and lifestyle data, returning a confidence score and full class-probability breakdown.
- Engineered derived features – BMI, lifestyle risk, age group, and city tier – directly in the Pydantic request schema using computed fields and validators, keeping the model input pipeline clean and stateless.

IMDB Sentiment Classification

Python, TensorFlow, Keras, NLP | [GitHub Link](#) | [Live Demo](#)

- Built a Recurrent Neural Network (Embedding → SimpleRNN → Dense/Sigmoid) for binary sentiment classification on 50K IMDB movie reviews, reaching ~82% validation accuracy.
- Handled sequence preprocessing (tokenization, padding, 10K-word vocabulary) and evaluated performance using accuracy and confusion-matrix analysis.

Customer Churn Prediction

Python, TensorFlow, Keras, Streamlit | [GitHub Link](#) | [Live Demo](#)

- Built and trained an ANN (Dense 64 → 32 → 1, ReLU/Sigmoid) on the Bank Customer Churn dataset with label encoding, one-hot encoding, and feature scaling, reaching ~86% validation accuracy using early stopping.
- Deployed as an interactive Streamlit app returning real-time churn probability scores from user-input feature values.

Achievements

- 500+ GitHub Commits
- GDG Solution Challenge– Google Developer Groups On Campus, 2025
- Smart India Hackathon 2025– Participant, Ministry of Education, Govt. of India, 2025

Certifications

- Claude with the Anthropic API — Anthropic, 2026
- Claude Code in Action — Anthropic, 2026
- Introduction to Agent Skills – Anthropic, 2026
- PyTorch for Deep Learning Bootcamp – Udemy 2026
- Introduction to Model Context Protocol (MCP) — Anthropic, 2026
- Machine Learning A-Z: AI, Python & R – Udemy (Nov 2025)